

# SPARK+AI SUMMIT 2020

---

Organized by  databricks®

# Text Extraction from Product Images


Rajesh Shreedhar Bhat

Data Scientist @Walmart Labs, Bengaluru

MS CS @ASU, Kaggle Competitions Expert

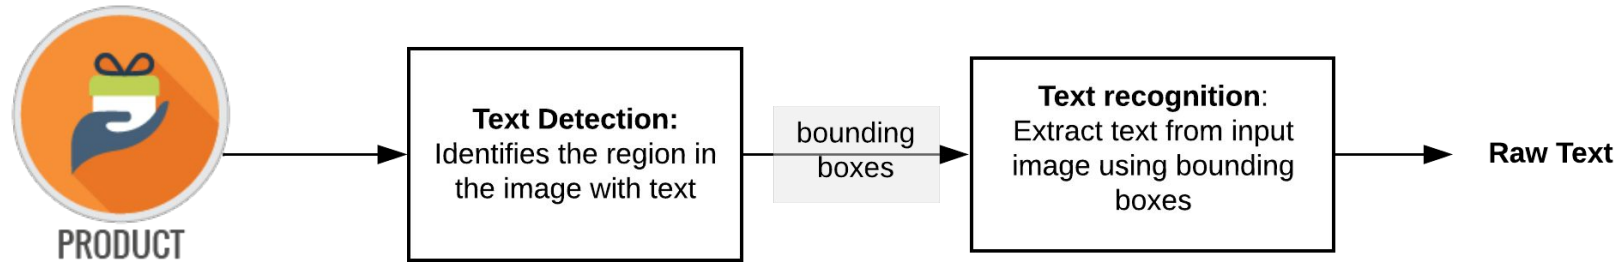
# Agenda

- Intro to Text Extraction
- **Text Detection**(TD)
- TD Model Architecture
- Training data generation
- **Text Recognition**(TR) training data preparation
- CRNN-CTC model for TR
- Receptive Fields
- CTC decoder and loss
- TR Training phase
- Other Advanced Techniques
- Questions ?







# Introduction : Text Extraction



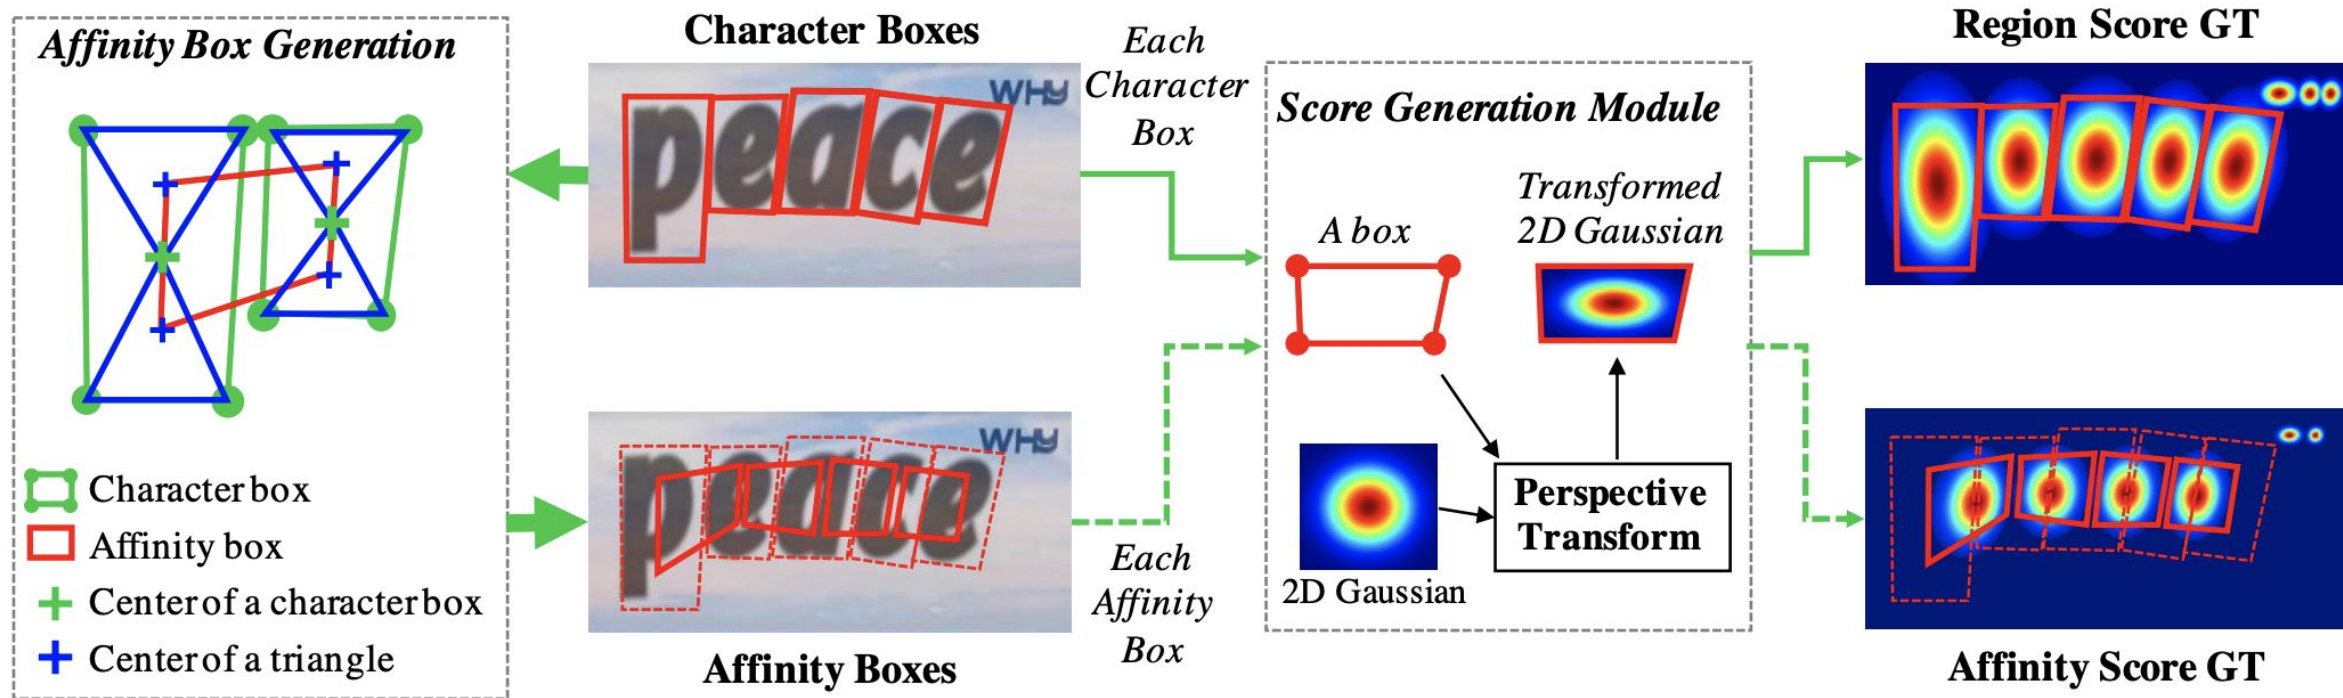
# Text Detection

# Image Segmentation – Input & Ground Truth





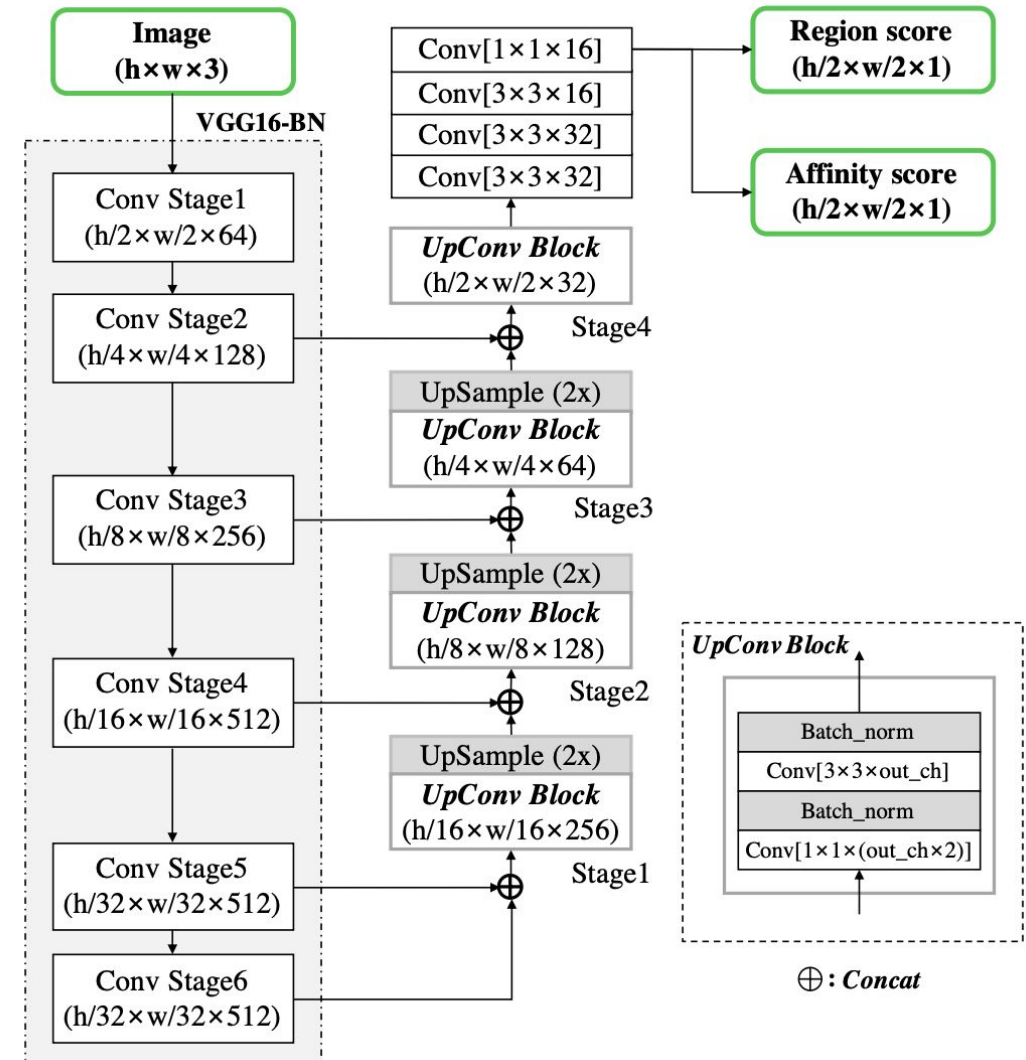
# Ground Truth Label Generation



# Text Detection – Model architecture


- VGG16 – BN as the backbone
- Model has skip connection in decoder part which is similar to U-Nets.
- Output :
  - Region score
  - Affinity score – grouping characters

**Ref:** Baek, Youngmin, et al. "Character Region Awareness for Text detection." *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 2019.







# Sample Output




Region Score



Affinity Score

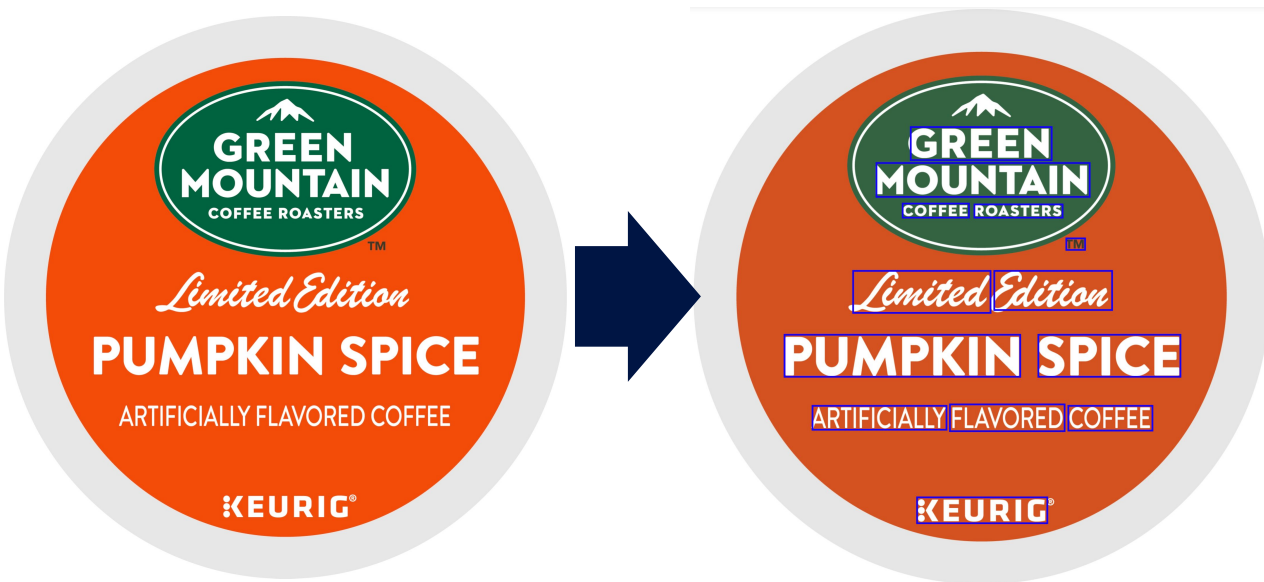


Region Score



Affinity Score

# Sample Output..





# Text Recognition



# Text Recognition – Training Data Preparation


**SynthText:** image generation engine for building a large annotated dataset.

**15 million** images generated with different **font styles, size, color & varying backgrounds** using product descriptions + open source datasets

**Vocabulary:** 92 characters  
Includes capital + small letters, numbers and special symbols




# Text Recognition CRNN CTC model



# CNN - Receptive Fields


- Receptive field is defined as the region in the input image/space that a particular CNN's feature is looking at.



two successive  
3x3 convolutions

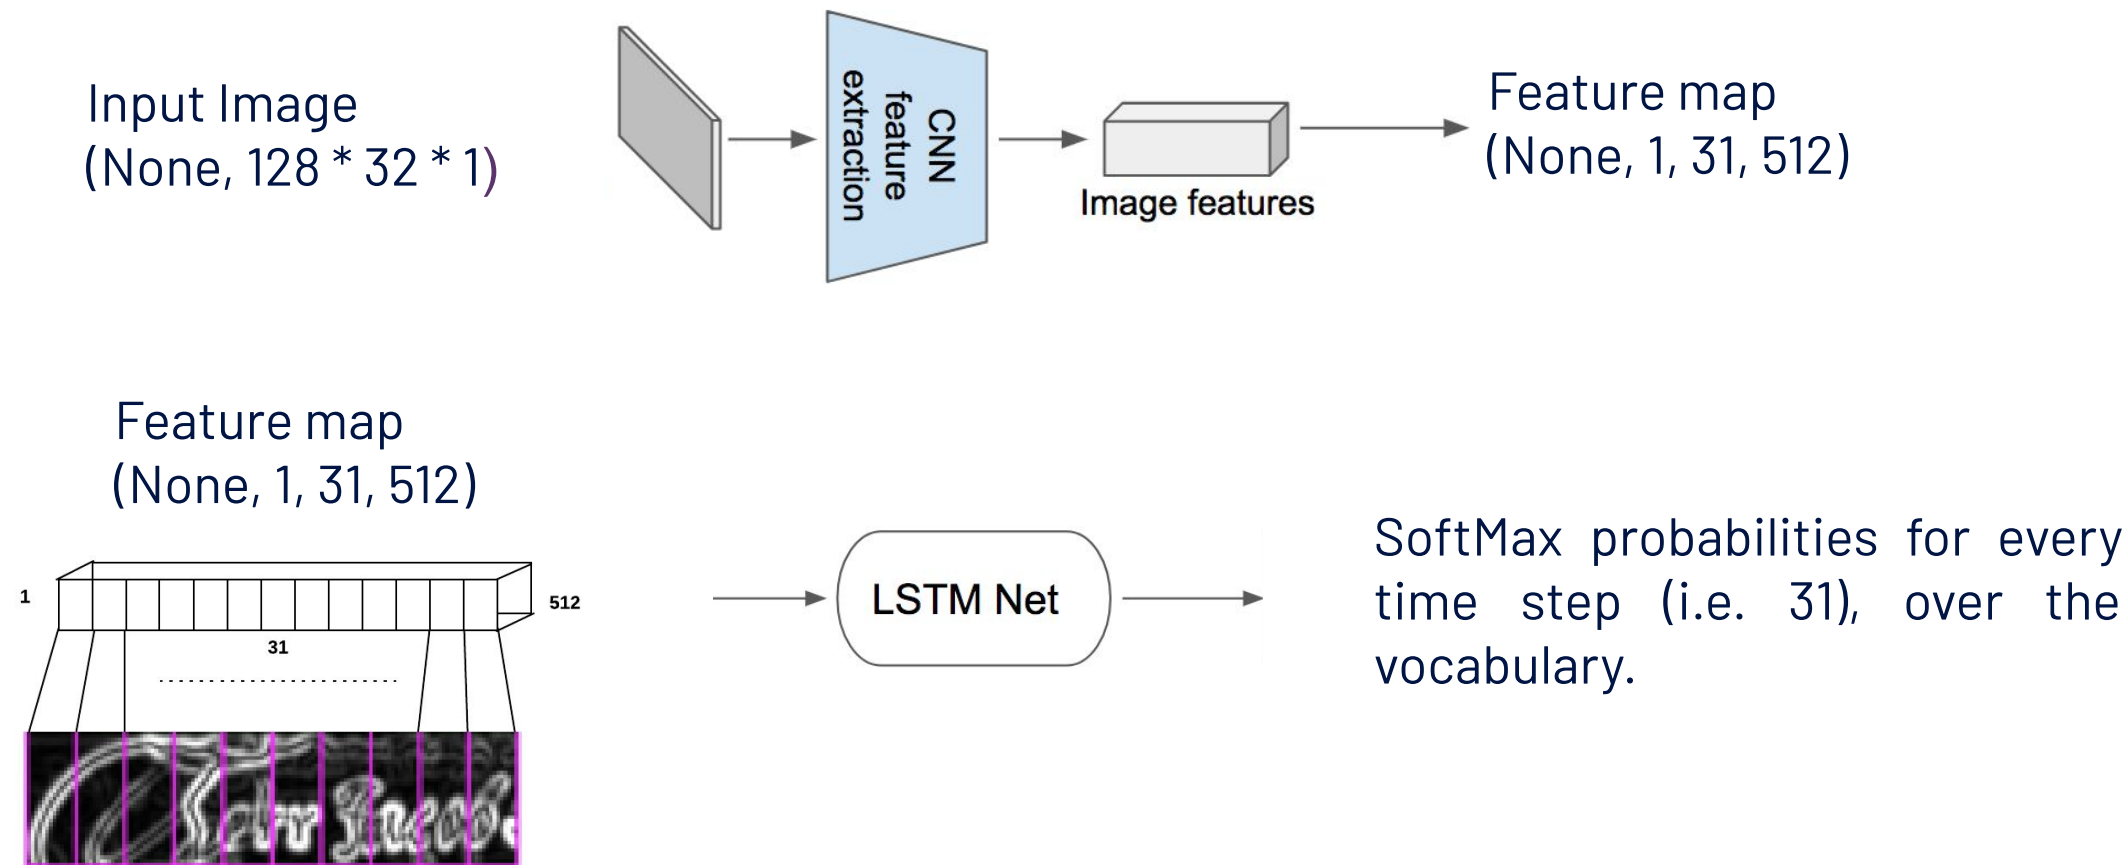
$$n_{out} = \left\lfloor \frac{n_{in} + 2p - k}{s} \right\rfloor + 1$$

Usage of intermediate layer features in SSD's in Object detection tasks.





# CNN features to LSTM




# Ground Truth and Output of TR task

Input Image




Ground Truth



Hello

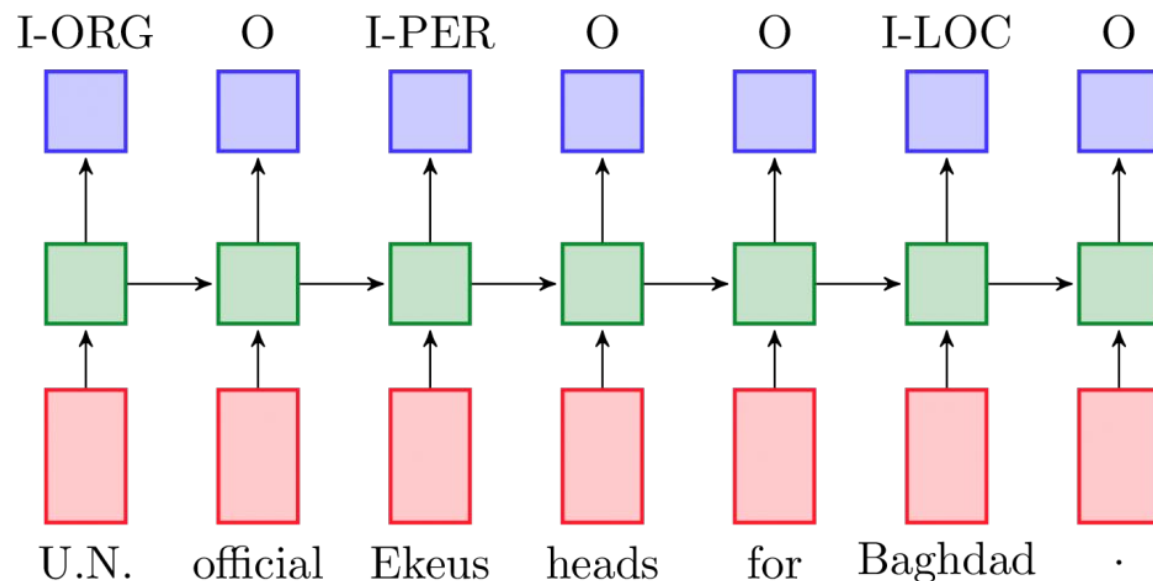
Output from LSTM model  
for 31 timesteps ..



Time step	t1	t2	t3	t4	t5	...	....	t27	t28	t29	t30	t31
Prediction	H	H	H	e	e	...	...	l	o	o	o	o

Length of Ground truth is **5** which is **not equal** to length of prediction i.e **31**

# How to calculate the loss?




NER model – loss: categorical cross entropy

- Do we have labels for every time steps of LSTM model in CRNN setting ?
- Can we use cross entropy loss?

**Answer is: NO!!**



# Mapping of Input to Output



Corresponding Text → Hello

The word 'Hello' written in a black, cursive script font.


Corresponding Text → Hello

Can we manually align each character to its location in the audio/image?

Yes!! But lot of manual effort is needed in creating training data.


# CTC to rescue

With just mapping from image to text and not worrying about alignment of each character to the location in input image, one should be able to train the network.



Merge repeats

Merge repeats  
Drop blank character



# Connectionist Temporal Connection (CTC) Loss

- Ground truth for an image **AB** → **AB**
- Vocabulary is { **A**, **B**, - }
- Let's say we have predictions for **3-time** steps from LSTM network (SoftMax probabilities over vocabulary at **t1, t2, t3**)
- Given that we use CTC decode operation discussed earlier, in which scenarios we can say output from the model is correct??



# CTC loss continued ..

**Ground Truth : AB**

t1	t2	t3
A	B	B
A	A	B
-	A	B
A	-	B
A	B	-

- Merge repeats

- Drop blank character

**AB**

SoftMax probabilities

	t1	t2	t3
A	0.8	0.7	0.1
B	0.1	0.1	0.8
-	0.1	0.2	0.1

Score for one path: **AAB** =  $(0.8 * 0.7 * 0.8)$  and similarly for other paths.

Probability of getting GT **AB**: =  $P(ABB) + P(AAB) + P(-AB) + P(A-B) + P(AB-)$

**Loss** :  $-\log(\text{Probability of getting GT})$

# CTC loss perfect match

Ground Truth : **A**

t1	t2	t3
A	-	-
-	A	-
-	-	A
-	A	A
A	A	-
A	-	A
A	A	A

- Merge repeats

- Drop blank character

**A**

- SoftMax probabilities

	t1	t2	t3
A	1	0	0
B	0	0	0
-	0	1	1

Score for one path: **A--** =  $(1 * 1 * 1)$  and similarly for other paths.

Probability of getting ground truth **A**: =  $P(A-) + P(-A-) + P(-A) + P(-AA) + P(AA-) + P(A-A) + P(AAA)$

**Loss** :  $-\log(\text{Probability of getting GT}) = 0$

# CTC loss perfect mismatch

**Ground Truth : A**

t1	t2	t3
A	-	-
-	A	-
-	-	A
-	A	A
A	A	-
A	-	A
A	A	A

- Merge repeats
- Drop blank character

**A**

SoftMax probabilities

	t1	t2	t3
A	0	0	0
B	1	1	1
-	0	0	0

Score for one path: **A - -** =  $(0 * 0 * 0)$  and similarly for other paths.

Probability of getting ground truth **A**: =  $P(A-) + P(-A-) + P(-A) + P(-AA) + P(AA-) + P(A-A) + P(AAA)$

**Loss** :  $-\log(\text{Probability of getting GT}) = \text{tends to infinity !!}$

# Model Architecture & CTC loss in TF

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 32, 128, 1)	0
conv2d_1 (Conv2D)	(None, 32, 128, 64)	640
max_pooling2d_1 (MaxPooling2)	(None, 16, 64, 64)	0
conv2d_2 (Conv2D)	(None, 16, 64, 128)	73856
max_pooling2d_2 (MaxPooling2)	(None, 8, 32, 128)	0
conv2d_3 (Conv2D)	(None, 8, 32, 256)	295168
conv2d_4 (Conv2D)	(None, 8, 32, 256)	590080
max_pooling2d_3 (MaxPooling2)	(None, 4, 32, 256)	0
conv2d_5 (Conv2D)	(None, 4, 32, 512)	1180160
batch_normalization_1 (Batch Normalization)	(None, 4, 32, 512)	2048
conv2d_6 (Conv2D)	(None, 4, 32, 512)	2359808
batch_normalization_2 (Batch Normalization)	(None, 4, 32, 512)	2048
max_pooling2d_4 (MaxPooling2)	(None, 2, 32, 512)	0
conv2d_7 (Conv2D)	(None, 1, 31, 512)	1049088
lambda_1 (Lambda)	(None, 31, 512)	0
bidirectional_1 (Bidirectional)	(None, 31, 256)	657408
bidirectional_2 (Bidirectional)	(None, 31, 256)	395264
dense_1 (Dense)	(None, 31, 93)	23901
Total params: 6,629,469		
Trainable params: 6,627,421		
Non-trainable params: 2,048		

```
tf.keras.backend.ctc_batch_cost(  
    y_true, y_pred, input_length, label_length  
)
```

## Arguments

<b>y_true</b>	tensor (samples, max_string_length) containing the truth labels.
<b>y_pred</b>	tensor (samples, time_steps, num_categories) containing the prediction, or output of the softmax.
<b>input_length</b>	tensor (samples, 1) containing the sequence length for each batch item in y_pred.
<b>label_length</b>	tensor (samples, 1) containing the sequence length for each batch item in y_true.

## Returns

Tensor with shape (samples,1) containing the CTC loss of each element.

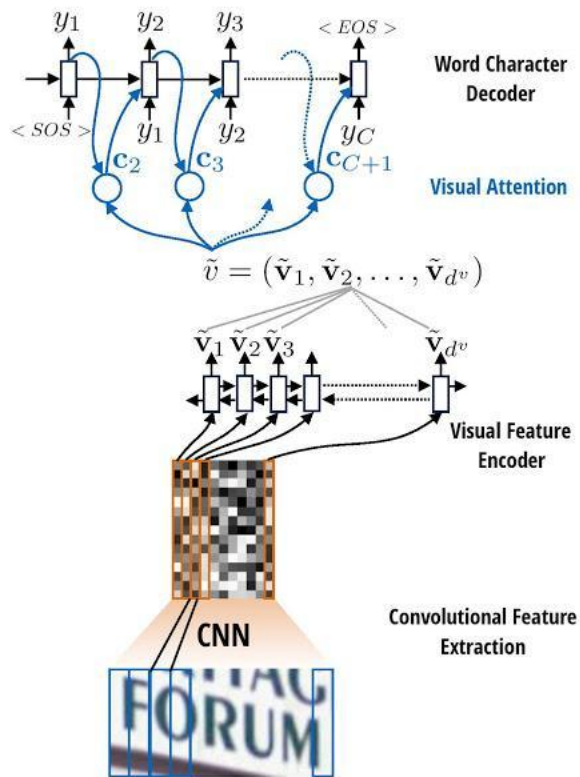
# Training Phase

- **15 million images ~ 690 GB when loaded into memory!!** Given that on an average images are of the shape **(128 \* 32 \* 3)** and **dtype is float32**.
- Usage of Python Generators to load only single batch in memory.
- Reducing the training time by using workers, max\_queue\_size & multi-processing in .fit\_generator in Keras.
- Training time ~ 2 hours for single epoch on single P100 GPU machine and prediction time ~1 sec for batch of 2048 images.

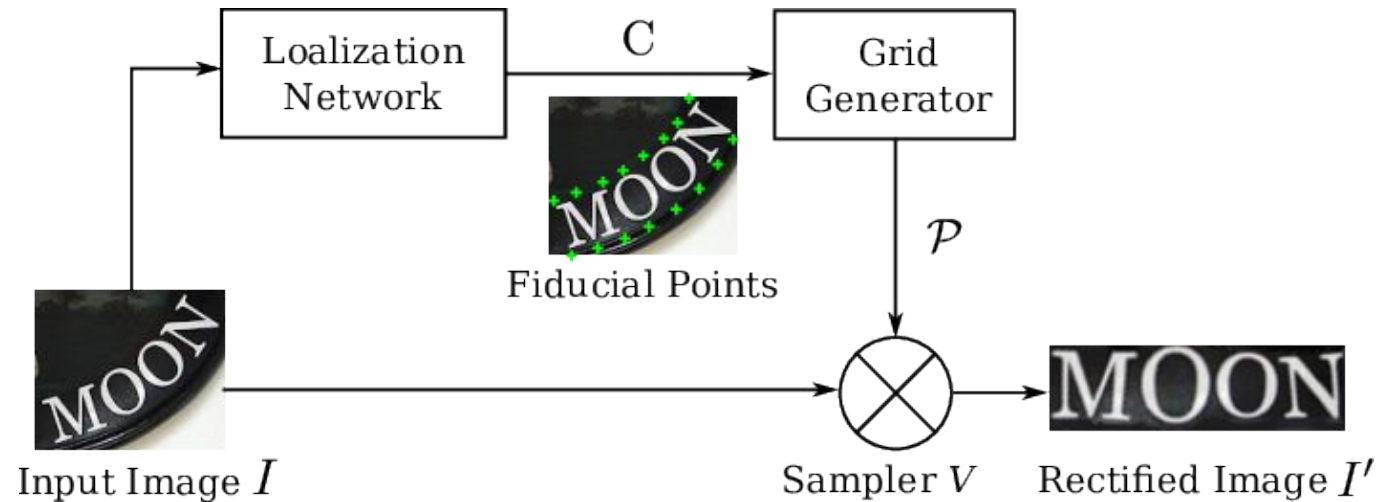


# Other Advanced Techniques

## Attention - OCR



## Spatial Transformer Network - before text recognition



**Ref:** [Jaderberg, Max, Karen Simonyan, and Andrew Zisserman. "Spatial transformer networks." Advances in neural information processing systems. 2015.](#)

# Code + PPT

<https://github.com/rajesh-bhat/spark-ai-summit-2020-text-extraction>



# Questions ??



*rsbhat@asu.edu*



<https://www.linkedin.com/in/rajeshshreedhar>